

# MP3\_P1A\_Introduction\_1.nbconvert

October 19, 2020

## 1 Assignment 3 Part 1A Introduction: Multi-label Image Classification

```
[1]: import os
import numpy as np
import torch
import torch.nn as nn
import torchvision

from torchvision import transforms
from sklearn.metrics import average_precision_score
from PIL import Image, ImageDraw
import matplotlib.pyplot as plt
from kaggle_submission import output_submission_csv
from classifier import SimpleClassifier, Classifier#, AlexNet
from voc_dataloader import VocDataset, VOC_CLASSES

%matplotlib inline
%load_ext autoreload
%autoreload 2
```

## 2 Multi-label Classification

In this assignment, you train a classifier to do multi-label classification on the PASCAL VOC 2007 dataset. The dataset has 20 different classes which can appear in any given image. Your classifier will predict whether each class appears in an image. This task is slightly different from exclusive multiclass classification like the ImageNet competition where only a single most appropriate class is predicted for an image.

### 2.1 Part 1A

You will use this notebook to warm up with pytorch and the code+dataset that we will use for assignment3.

#### 2.1.1 What to do

In part 1A, You are asked to run below experiments. You don't need to change hyperparameters for this Part 1A's experiments. (the following code provides everything that you will need.) 1. to

train a simple network (defined in `classifiers.py`) 2. to train the AlexNet (PyTorch built-in) - from scratch - finetuning AlexNet pretrained on ImageNet

### 2.1.2 What to submit

We ask you to run the following code and report the results in your homework submission. You may want to leverage this part 1A get yourself familiar with PyTorch.

You will need the numbers and plots this notebook outputs for reports, but you are not required to submit this notebook as a printed pdf.

## 2.2 Reading Pascal Data

### 2.2.1 Loading Training Data

In the following cell we will load the training data and also apply some transforms to the data.

```
[2]: # Transforms applied to the training data
normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std= [0.229, 0.224, 0.225])

train_transform = transforms.Compose([
    transforms.Resize(227),
    transforms.CenterCrop(227),
    transforms.ToTensor(),
    normalize
])

[3]: ds_train = VocDataset('VOCdevkit_2007/VOC2007/', 'train', train_transform)
```

### 2.2.2 Loading Validation Data

We will load the test data for the PASCAL VOC 2007 dataset. Do **NOT** add data augmentation transforms to validation data.

```
[4]: # Transforms applied to the testing data
test_transform = transforms.Compose([
    transforms.Resize(227),
    transforms.CenterCrop(227),
    transforms.ToTensor(),
    normalize,
])

[5]: ds_val = VocDataset('VOCdevkit_2007/VOC2007/', 'val', test_transform)
```


### 2.2.3 Visualizing the Data

PASCAL VOC has bounding box annotations in addition to class labels. Use the following code to visualize some random examples and corresponding annotations from the train set.

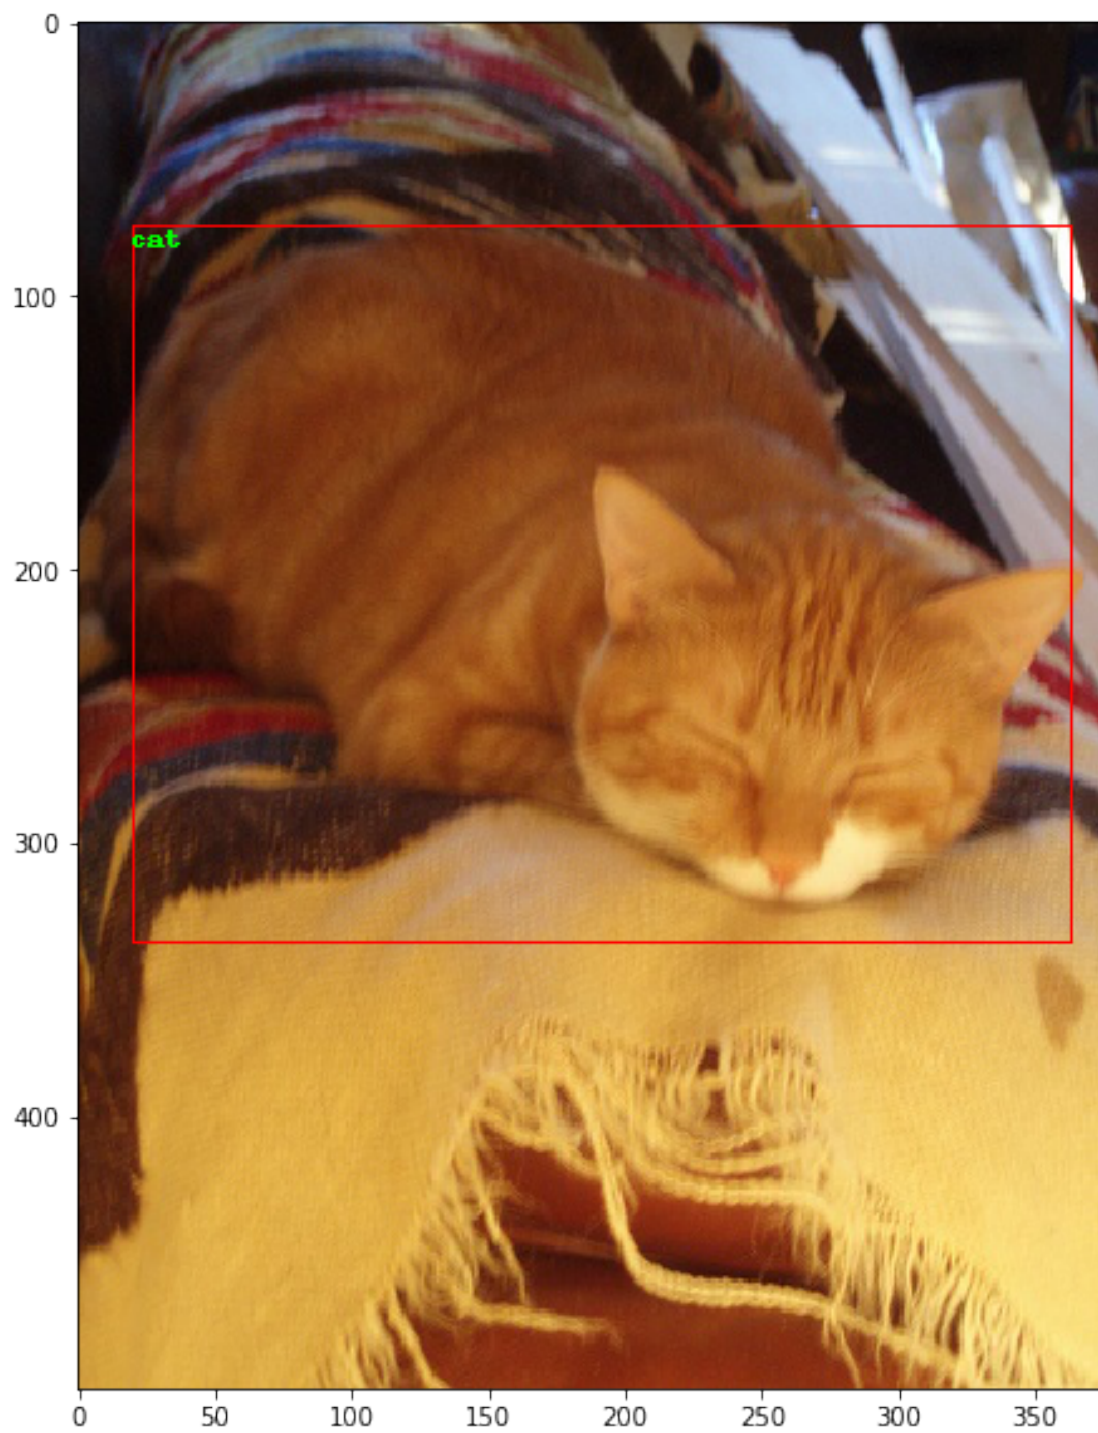
```
[6]: for i in range(5):
      idx = np.random.randint(0, len(ds_train.names)+1)
      _imgpath = os.path.join('VOCdevkit_2007/VOC2007/', 'JPEGImages', ds_train.
      ↪names[idx]+' .jpg')
      img = Image.open(_imgpath).convert('RGB')
      draw = ImageDraw.Draw(img)
      for j in range(len(ds_train.box_indices[idx])):
          obj = ds_train.box_indices[idx][j]
          draw.rectangle(list(obj), outline=(255,0,0))
          draw.text(list(obj[0:2]), ds_train.classes[ds_train.
          ↪label_order[idx][j]], fill=(0,255,0))
      plt.figure(figsize = (10,10))
      plt.imshow(np.array(img))
```














### 3 Classification

```
[7]: # declare what device to use: gpu/cpu
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda:0

```
[8]: train_loader = torch.utils.data.DataLoader(dataset=ds_train,
                                                batch_size=50,
                                                shuffle=True,
                                                num_workers=1)
```

```
[9]: val_loader = torch.utils.data.DataLoader(dataset=ds_val,
                                                batch_size=50,
                                                shuffle=True,
                                                num_workers=1)
```

```
[10]: def train_classifier(train_loader, classifier, criterion, optimizer):
    classifier.train()
    loss_ = 0.0
    losses = []
    for i, (images, labels, _) in enumerate(train_loader):
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        logits = classifier(images)
        loss = criterion(logits, labels)
        loss.backward()
        optimizer.step()
        losses.append(loss)
    return torch.stack(losses).mean().item()

[11]: def test_classifier(test_loader, classifier, criterion, print_ind_classes=True,
    ↪print_total=True):
    classifier.eval()
    losses = []
    with torch.no_grad():
        y_true = np.zeros((0,21))
        y_score = np.zeros((0,21))
        for i, (images, labels, _) in enumerate(test_loader):
            images, labels = images.to(device), labels.to(device)
            logits = classifier(images)
            y_true = np.concatenate((y_true, labels.cpu().numpy()), axis=0)
            y_score = np.concatenate((y_score, logits.cpu().numpy()), axis=0)
            loss = criterion(logits, labels)
            losses.append(loss.item())
    aps = []
    # ignore first class which is background
    for i in range(1, y_true.shape[1]):
        ap = average_precision_score(y_true[:, i], y_score[:, i])
        if print_ind_classes:
            print('----- Class: {:<12}      AP: {:>8.4f} -----'.
    ↪format(VOC_CLASSES[i], ap))
        aps.append(ap)

    mAP = np.mean(aps)
    test_loss = np.mean(losses)
    if print_total:
        print('mAP: {0:.4f}'.format(mAP))
        print('Avg loss: {}'.format(test_loss))

    return mAP, test_loss, aps

[12]: # plot functions
def plot_losses(train, val, test_frequency, num_epochs):
```



```

plt.plot(train, label="train")
indices = [i for i in range(num_epochs) if ((i+1)%test_frequency == 0 or i
↪==0)]
plt.plot(indices, val, label="val")
plt.title("Loss Plot")
plt.ylabel("Loss")
plt.xlabel("Epoch")
plt.legend()
plt.show()

def plot_mAP(train, val, test_frequency, num_epochs):
    indices = [i for i in range(num_epochs) if ((i+1)%test_frequency == 0 or i
↪==0)]
    plt.plot(indices, train, label="train")
    plt.plot(indices, val, label="val")
    plt.title("mAP Plot")
    plt.ylabel("mAP")
    plt.xlabel("Epoch")
    plt.legend()
    plt.show()

```

### 3.1 Training the network

The simple network you are given as is will allow you to reach around 0.15-0.2 mAP. In this project, you will find ways to design a better network. Save plots and final test mAP scores as you will be adding these to the writeup.

```

[13]: def train(classifier, num_epochs, train_loader, val_loader, criterion,
↪optimizer, test_frequency=5):
    train_losses = []
    train_mAPs = []
    val_losses = []
    val_mAPs = []

    for epoch in range(1,num_epochs+1):
        print("Starting epoch number " + str(epoch))
        train_loss = train_classifier(train_loader, classifier, criterion,
↪optimizer)
        train_losses.append(train_loss)
        print("Loss for Training on Epoch " +str(epoch) + " is "+
↪str(train_loss))
        if(epoch%test_frequency==0 or epoch==1):
            mAP_train, _, _ = test_classifier(train_loader, classifier,
↪criterion, False, False)
            train_mAPs.append(mAP_train)

```

```

        mAP_val, val_loss, _ = test_classifier(val_loader, classifier, ↵
↵criterion)
        print('Evaluating classifier')
        print("Mean Precision Score for Testing on Epoch " +str(epoch) + "↵
↵is "+ str(mAP_val))
        val_losses.append(val_loss)
        val_mAPs.append(mAP_val)

    return classifier, train_losses, val_losses, train_mAPs, val_mAPs

```

```

[14]: classifier = SimpleClassifier().to(device)
      # You can use this function to reload a network you have already saved ↵
      ↵previously
      #classifier.load_state_dict(torch.load('voc_classifier.pth'))

```

```

[15]: criterion = nn.MultiLabelSoftMarginLoss()
      optimizer = torch.optim.SGD(classifier.parameters(), lr=0.01, momentum=0.9)

```

```

[16]: # Training the Classifier
      num_epochs = 20
      test_frequency = 5

      classifier, train_losses, val_losses, train_mAPs, val_mAPs = train(classifier, ↵
      ↵num_epochs, train_loader, val_loader, criterion, optimizer, test_frequency)

```

Starting epoch number 1

Loss for Training on Epoch 1 is 0.5439279079437256

-----	Class: aeroplane	AP:	0.1492	-----
-----	Class: bicycle	AP:	0.0513	-----
-----	Class: bird	AP:	0.1055	-----
-----	Class: boat	AP:	0.1174	-----
-----	Class: bottle	AP:	0.0393	-----
-----	Class: bus	AP:	0.0316	-----
-----	Class: car	AP:	0.1154	-----
-----	Class: cat	AP:	0.0791	-----
-----	Class: chair	AP:	0.1124	-----
-----	Class: cow	AP:	0.0298	-----
-----	Class: diningtable	AP:	0.0470	-----
-----	Class: dog	AP:	0.1108	-----
-----	Class: horse	AP:	0.0495	-----
-----	Class: motorbike	AP:	0.0468	-----
-----	Class: person	AP:	0.4366	-----
-----	Class: pottedplant	AP:	0.0443	-----
-----	Class: sheep	AP:	0.0237	-----
-----	Class: sofa	AP:	0.0838	-----
-----	Class: train	AP:	0.0363	-----
-----	Class: tvmonitor	AP:	0.0502	-----

mAP: 0.0880  
 Avg loss: 0.27371764679749805  
 Evaluating classifier  
 Mean Precision Score for Testing on Epoch 1 is 0.0880146134797726  
 Starting epoch number 2  
 Loss for Training on Epoch 2 is 0.24799811840057373  
 Starting epoch number 3  
 Loss for Training on Epoch 3 is 0.24086320400238037  
 Starting epoch number 4  
 Loss for Training on Epoch 4 is 0.2410968840122223  
 Starting epoch number 5  
 Loss for Training on Epoch 5 is 0.24504314363002777  

-----	Class: aeroplane	AP: 0.0861	-----
-----	Class: bicycle	AP: 0.0455	-----
-----	Class: bird	AP: 0.1038	-----
-----	Class: boat	AP: 0.0756	-----
-----	Class: bottle	AP: 0.0446	-----
-----	Class: bus	AP: 0.0284	-----
-----	Class: car	AP: 0.1113	-----
-----	Class: cat	AP: 0.1066	-----
-----	Class: chair	AP: 0.1316	-----
-----	Class: cow	AP: 0.0305	-----
-----	Class: diningtable	AP: 0.0553	-----
-----	Class: dog	AP: 0.1298	-----
-----	Class: horse	AP: 0.0486	-----
-----	Class: motorbike	AP: 0.0414	-----
-----	Class: person	AP: 0.3543	-----
-----	Class: pottedplant	AP: 0.0466	-----
-----	Class: sheep	AP: 0.0263	-----
-----	Class: sofa	AP: 0.1045	-----
-----	Class: train	AP: 0.0353	-----
-----	Class: tvmonitor	AP: 0.0539	-----

mAP: 0.0830  
 Avg loss: 0.23721890998821632  
 Evaluating classifier  
 Mean Precision Score for Testing on Epoch 5 is 0.08300479102381039  
 Starting epoch number 6  
 Loss for Training on Epoch 6 is 0.24111507833003998  
 Starting epoch number 7  
 Loss for Training on Epoch 7 is 0.2412501573562622  
 Starting epoch number 8  
 Loss for Training on Epoch 8 is 0.23983041942119598  
 Starting epoch number 9  
 Loss for Training on Epoch 9 is 0.2369721233844757  
 Starting epoch number 10  
 Loss for Training on Epoch 10 is 0.2390352338552475  

-----	Class: aeroplane	AP: 0.0745	-----
-----	Class: bicycle	AP: 0.0440	-----

-----	Class: bird	AP: 0.1023	-----
-----	Class: boat	AP: 0.0600	-----
-----	Class: bottle	AP: 0.0456	-----
-----	Class: bus	AP: 0.0278	-----
-----	Class: car	AP: 0.1143	-----
-----	Class: cat	AP: 0.1177	-----
-----	Class: chair	AP: 0.1397	-----
-----	Class: cow	AP: 0.0315	-----
-----	Class: diningtable	AP: 0.0585	-----
-----	Class: dog	AP: 0.1352	-----
-----	Class: horse	AP: 0.0501	-----
-----	Class: motorbike	AP: 0.0411	-----
-----	Class: person	AP: 0.3610	-----
-----	Class: pottedplant	AP: 0.0480	-----
-----	Class: sheep	AP: 0.0297	-----
-----	Class: sofa	AP: 0.1106	-----
-----	Class: train	AP: 0.0358	-----
-----	Class: tvmonitor	AP: 0.0560	-----

mAP: 0.0842

Avg loss: 0.23516886199221892

Evaluating classifier

Mean Precision Score for Testing on Epoch 10 is 0.08417308886982525

Starting epoch number 11

Loss for Training on Epoch 11 is 0.24059483408927917

Starting epoch number 12

Loss for Training on Epoch 12 is 0.23849020898342133

Starting epoch number 13

Loss for Training on Epoch 13 is 0.23976543545722961

Starting epoch number 14

Loss for Training on Epoch 14 is 0.2394586205482483

Starting epoch number 15

Loss for Training on Epoch 15 is 0.23732246458530426

-----	Class: aeroplane	AP: 0.0859	-----
-----	Class: bicycle	AP: 0.0446	-----
-----	Class: bird	AP: 0.1060	-----
-----	Class: boat	AP: 0.0545	-----
-----	Class: bottle	AP: 0.0467	-----
-----	Class: bus	AP: 0.0285	-----
-----	Class: car	AP: 0.1269	-----
-----	Class: cat	AP: 0.1185	-----
-----	Class: chair	AP: 0.1529	-----
-----	Class: cow	AP: 0.0331	-----
-----	Class: diningtable	AP: 0.0633	-----
-----	Class: dog	AP: 0.1312	-----
-----	Class: horse	AP: 0.0531	-----
-----	Class: motorbike	AP: 0.0424	-----
-----	Class: person	AP: 0.3945	-----
-----	Class: pottedplant	AP: 0.0508	-----



```



----- Class: sheep          AP:  0.0410 -----
----- Class: sofa           AP:  0.1130 -----
----- Class: train          AP:  0.0385 -----
----- Class: tvmonitor      AP:  0.0568 -----
mAP: 0.0891
Avg loss: 0.23332357172872506
Evaluating classifier
Mean Precision Score for Testing on Epoch 15 is 0.08911148094568158
Starting epoch number 16
Loss for Training on Epoch 16 is 0.2368842512369156
Starting epoch number 17
Loss for Training on Epoch 17 is 0.23557999730110168
Starting epoch number 18
Loss for Training on Epoch 18 is 0.23437541723251343
Starting epoch number 19
Loss for Training on Epoch 19 is 0.23336078226566315
Starting epoch number 20
Loss for Training on Epoch 20 is 0.23432353138923645
----- Class: aeroplane      AP:  0.2138 -----
----- Class: bicycle        AP:  0.0576 -----
----- Class: bird           AP:  0.1294 -----
----- Class: boat           AP:  0.1173 -----
----- Class: bottle         AP:  0.0618 -----
----- Class: bus            AP:  0.0552 -----
----- Class: car            AP:  0.1758 -----
----- Class: cat            AP:  0.1182 -----
----- Class: chair          AP:  0.1766 -----
----- Class: cow            AP:  0.0392 -----
----- Class: diningtable    AP:  0.0909 -----
----- Class: dog            AP:  0.1269 -----
----- Class: horse          AP:  0.0603 -----
----- Class: motorbike      AP:  0.0464 -----
----- Class: person         AP:  0.4821 -----
----- Class: pottedplant     AP:  0.0538 -----
----- Class: sheep          AP:  0.0408 -----
----- Class: sofa           AP:  0.1204 -----
----- Class: train          AP:  0.0926 -----
----- Class: tvmonitor      AP:  0.0572 -----
mAP: 0.1158
Avg loss: 0.22855493046489417
Evaluating classifier
Mean Precision Score for Testing on Epoch 20 is 0.11581629322238615

```

```

[17]: # Compare train and validation metrics
      plot_losses(train_losses, val_losses, test_frequency, num_epochs)
      plot_mAP(train_mAPs, val_mAPs, test_frequency, num_epochs)

```



```
[18]: # Save the classifier network
# Suggestion: you can save checkpoints of your network during training and
# reload them later
torch.save(classifier.state_dict(), './voc_simple_classifier.pth')
```

## 4 Evaluate on test set

```
[19]: ds_test = VocDataset('VOCdevkit_2007/VOC2007test/', 'test', test_transform)

test_loader = torch.utils.data.DataLoader(dataset=ds_test,
                                           batch_size=50,
                                           shuffle=False,
                                           num_workers=1)

# Transforms applied to the testing data
test_transform = transforms.Compose([
    transforms.Resize(227),
    transforms.CenterCrop(227),
    transforms.ToTensor(),
    normalize,
])

mAP_test, test_loss, test_aps = test_classifier(test_loader, classifier,
# criterion)
```

```
----- Class: aeroplane      AP:  0.1757 -----
----- Class: bicycle        AP:  0.0535 -----
----- Class: bird           AP:  0.1020 -----
----- Class: boat           AP:  0.1069 -----
----- Class: bottle         AP:  0.0565 -----
----- Class: bus            AP:  0.0493 -----
----- Class: car            AP:  0.1842 -----
----- Class: cat            AP:  0.1158 -----
----- Class: chair          AP:  0.1562 -----
----- Class: cow            AP:  0.0473 -----
----- Class: diningtable    AP:  0.0665 -----
----- Class: dog            AP:  0.1348 -----
----- Class: horse          AP:  0.0573 -----
----- Class: motorbike      AP:  0.0455 -----
----- Class: person         AP:  0.4929 -----
----- Class: pottedplant    AP:  0.0504 -----
----- Class: sheep          AP:  0.0595 -----
----- Class: sofa           AP:  0.1157 -----
----- Class: train          AP:  0.1170 -----
----- Class: tvmonitor      AP:  0.0565 -----
mAP: 0.1122
```

Avg loss: 0.2248046140372753

```
[20]: output_submission_csv('my_solution.csv', test_aps)
```

## 5 AlexNet Baselines (From Scratch)

AlexNet was one of the earliest deep learning models to have success in classification. In this section we will be running classification with AlexNet as a baseline. Furthermore, we will run an ImageNet-pretrained AlexNet to observe the impact of well-trained features. Save plots and final test mAP scores as you will be adding these to the writeup.

### 5.1 Running AlexNet

In this section, we train AlexNet from scratch using the same hyperparameters as our previous experiment.

```
[21]: num_epochs = 20
      test_frequency = 5

      # Change classifier to AlexNet
      classifier = torchvision.models.alexnet(pretrained=False)
      classifier.classifier._modules['6'] = nn.Linear(4096, 21)
      classifier = classifier.to(device)

      criterion = nn.MultiLabelSoftMarginLoss()

      optimizer = torch.optim.SGD(classifier.parameters(), lr=0.01, momentum=0.9)

[22]: classifier, train_losses, val_losses, train_mAPs, val_mAPs = train(classifier,
      ↪ num_epochs, train_loader, val_loader, criterion, optimizer, test_frequency)
```

Starting epoch number 1

Loss for Training on Epoch 1 is 0.5570489168167114

----- Class: aeroplane	AP: 0.1382	-----
----- Class: bicycle	AP: 0.0430	-----
----- Class: bird	AP: 0.1053	-----
----- Class: boat	AP: 0.1274	-----
----- Class: bottle	AP: 0.0376	-----
----- Class: bus	AP: 0.0292	-----
----- Class: car	AP: 0.1091	-----
----- Class: cat	AP: 0.0995	-----
----- Class: chair	AP: 0.1093	-----
----- Class: cow	AP: 0.0361	-----
----- Class: diningtable	AP: 0.0458	-----
----- Class: dog	AP: 0.1130	-----
----- Class: horse	AP: 0.0451	-----
----- Class: motorbike	AP: 0.0339	-----
----- Class: person	AP: 0.3249	-----



```

----- Class: pottedplant      AP:  0.0409 -----
----- Class: sheep            AP:  0.0379 -----
----- Class: sofa             AP:  0.0885 -----
----- Class: train            AP:  0.0405 -----
----- Class: tvmonitor        AP:  0.0473 -----
mAP: 0.0826
Avg loss: 0.30657491087913513
Evaluating classifier
Mean Precision Score for Testing on Epoch 1 is 0.08262655529072144
Starting epoch number 2
Loss for Training on Epoch 2 is 0.24965687096118927
Starting epoch number 3
Loss for Training on Epoch 3 is 0.24649912118911743
Starting epoch number 4
Loss for Training on Epoch 4 is 0.2476574331521988
Starting epoch number 5
Loss for Training on Epoch 5 is 0.24382369220256805
----- Class: aeroplane        AP:  0.0945 -----
----- Class: bicycle          AP:  0.0463 -----
----- Class: bird             AP:  0.1172 -----
----- Class: boat             AP:  0.1018 -----
----- Class: bottle           AP:  0.0414 -----
----- Class: bus              AP:  0.0292 -----
----- Class: car              AP:  0.1124 -----
----- Class: cat              AP:  0.0929 -----
----- Class: chair            AP:  0.1203 -----
----- Class: cow              AP:  0.0358 -----
----- Class: diningtable      AP:  0.0527 -----
----- Class: dog              AP:  0.1176 -----
----- Class: horse            AP:  0.0485 -----
----- Class: motorbike        AP:  0.0371 -----
----- Class: person           AP:  0.3513 -----
----- Class: pottedplant      AP:  0.0459 -----
----- Class: sheep            AP:  0.0414 -----
----- Class: sofa             AP:  0.0938 -----
----- Class: train            AP:  0.0403 -----
----- Class: tvmonitor        AP:  0.0480 -----
mAP: 0.0834
Avg loss: 0.243446968349756
Evaluating classifier
Mean Precision Score for Testing on Epoch 5 is 0.08341850320114871
Starting epoch number 6
Loss for Training on Epoch 6 is 0.2430320680141449
Starting epoch number 7
Loss for Training on Epoch 7 is 0.2422732263803482
Starting epoch number 8
Loss for Training on Epoch 8 is 0.24190865457057953
Starting epoch number 9

```

Loss for Training on Epoch 9 is 0.24298124015331268

Starting epoch number 10

Loss for Training on Epoch 10 is 0.2408696413040161

-----	Class: aeroplane	AP:	0.0836	-----
-----	Class: bicycle	AP:	0.0487	-----
-----	Class: bird	AP:	0.1243	-----
-----	Class: boat	AP:	0.0913	-----
-----	Class: bottle	AP:	0.0413	-----
-----	Class: bus	AP:	0.0298	-----
-----	Class: car	AP:	0.1171	-----
-----	Class: cat	AP:	0.0922	-----
-----	Class: chair	AP:	0.1276	-----
-----	Class: cow	AP:	0.0365	-----
-----	Class: diningtable	AP:	0.0546	-----
-----	Class: dog	AP:	0.1200	-----
-----	Class: horse	AP:	0.0515	-----
-----	Class: motorbike	AP:	0.0395	-----
-----	Class: person	AP:	0.4411	-----
-----	Class: pottedplant	AP:	0.0494	-----
-----	Class: sheep	AP:	0.0475	-----
-----	Class: sofa	AP:	0.0935	-----
-----	Class: train	AP:	0.0412	-----
-----	Class: tvmonitor	AP:	0.0496	-----

mAP: 0.0890

Avg loss: 0.23791137483774447

Evaluating classifier

Mean Precision Score for Testing on Epoch 10 is 0.08900571633952556

Starting epoch number 11

Loss for Training on Epoch 11 is 0.24085105955600739

Starting epoch number 12

Loss for Training on Epoch 12 is 0.24197809398174286

Starting epoch number 13

Loss for Training on Epoch 13 is 0.24029192328453064

Starting epoch number 14

Loss for Training on Epoch 14 is 0.2365572601556778

Starting epoch number 15

Loss for Training on Epoch 15 is 0.24076105654239655

-----	Class: aeroplane	AP:	0.0848	-----
-----	Class: bicycle	AP:	0.0520	-----
-----	Class: bird	AP:	0.1327	-----
-----	Class: boat	AP:	0.0803	-----
-----	Class: bottle	AP:	0.0438	-----
-----	Class: bus	AP:	0.0312	-----
-----	Class: car	AP:	0.1278	-----
-----	Class: cat	AP:	0.0920	-----
-----	Class: chair	AP:	0.1425	-----
-----	Class: cow	AP:	0.0373	-----
-----	Class: diningtable	AP:	0.0606	-----

-----	Class: dog	AP:	0.1209	-----
-----	Class: horse	AP:	0.0549	-----
-----	Class: motorbike	AP:	0.0435	-----
-----	Class: person	AP:	0.4484	-----
-----	Class: pottedplant	AP:	0.0547	-----
-----	Class: sheep	AP:	0.0502	-----
-----	Class: sofa	AP:	0.0982	-----
-----	Class: train	AP:	0.0438	-----
-----	Class: tvmonitor	AP:	0.0518	-----

mAP: 0.0926

Avg loss: 0.2371828433345346

Evaluating classifier

Mean Precision Score for Testing on Epoch 15 is 0.092551745332285

Starting epoch number 16

Loss for Training on Epoch 16 is 0.2388356626033783

Starting epoch number 17

Loss for Training on Epoch 17 is 0.2375527024269104

Starting epoch number 18

Loss for Training on Epoch 18 is 0.2359590083360672

Starting epoch number 19

Loss for Training on Epoch 19 is 0.23879577219486237

Starting epoch number 20

Loss for Training on Epoch 20 is 0.23340480029582977

-----	Class: aeroplane	AP:	0.2291	-----
-----	Class: bicycle	AP:	0.0613	-----
-----	Class: bird	AP:	0.1364	-----
-----	Class: boat	AP:	0.1809	-----
-----	Class: bottle	AP:	0.0589	-----
-----	Class: bus	AP:	0.0441	-----
-----	Class: car	AP:	0.1714	-----
-----	Class: cat	AP:	0.1006	-----
-----	Class: chair	AP:	0.1746	-----
-----	Class: cow	AP:	0.0399	-----
-----	Class: diningtable	AP:	0.0879	-----
-----	Class: dog	AP:	0.1312	-----
-----	Class: horse	AP:	0.0580	-----
-----	Class: motorbike	AP:	0.0465	-----
-----	Class: person	AP:	0.5024	-----
-----	Class: pottedplant	AP:	0.0672	-----
-----	Class: sheep	AP:	0.0485	-----
-----	Class: sofa	AP:	0.1202	-----
-----	Class: train	AP:	0.0756	-----
-----	Class: tvmonitor	AP:	0.0536	-----



mAP: 0.1194

Avg loss: 0.23350198099426195

Evaluating classifier

Mean Precision Score for Testing on Epoch 20 is 0.11941861495429547

```
[23]: plot_losses(train_losses, val_losses, test_frequency, num_epochs)
      plot_mAP(train_mAPs, val_mAPs, test_frequency, num_epochs)
```





```
[24]: mAP_test, test_loss, test_aps = test_classifier(test_loader, classifier,
↪ criterion)
print("Test mAP: ", mAP_test)
```

```
----- Class: aeroplane      AP:  0.2001  -----
----- Class: bicycle       AP:  0.0579  -----
----- Class: bird          AP:  0.1106  -----
----- Class: boat          AP:  0.1395  -----
----- Class: bottle        AP:  0.0568  -----
----- Class: bus           AP:  0.0414  -----
----- Class: car           AP:  0.1696  -----
----- Class: cat           AP:  0.1009  -----
----- Class: chair         AP:  0.1592  -----
----- Class: cow           AP:  0.0478  -----
----- Class: diningtable   AP:  0.0700  -----
----- Class: dog           AP:  0.1367  -----
----- Class: horse         AP:  0.0553  -----
----- Class: motorbike     AP:  0.0417  -----
----- Class: person        AP:  0.5110  -----
----- Class: pottedplant   AP:  0.0569  -----
----- Class: sheep         AP:  0.0651  -----
----- Class: sofa          AP:  0.1209  -----
----- Class: train         AP:  0.0885  -----
----- Class: tvmonitor     AP:  0.0537  -----
mAP: 0.1142
Avg loss: 0.23103498727083205
Test mAP: 0.11418412719617296
```

You should notice somewhat poor performance. You could try running AlexNet with an Adam optimizer instead with learning rate  $1e-4$  to see if that makes a difference. This experiment is not required for the writeup, but it may show you the importance of a good learning rate and optimizer.

## 5.2 Pretrained AlexNet

Here we look at the impact of pretrained features. This model's weights were trained on ImageNet, which is a much larger dataset. How do pretrained features perform on VOC? Why do you think there is such a large difference in performance?

```
[25]: num_epochs = 20
test_frequency = 5

# Load Pretrained AlexNet
classifier = torchvision.models.alexnet(pretrained=True)
classifier.classifier._modules['6'] = nn.Linear(4096, 21)
classifier = classifier.to(device)
optimizer = torch.optim.SGD(classifier.parameters(), lr=0.01, momentum=0.9)
```

```
[26]: classifier, train_losses, val_losses, train_mAPs, val_mAPs = train(classifier,
    ↪ num_epochs, train_loader, val_loader, criterion, optimizer, test_frequency)
```

```
Starting epoch number 1
Loss for Training on Epoch 1 is 0.2252361923456192
----- Class: aeroplane      AP:   0.7824 -----
----- Class: bicycle        AP:   0.5411 -----
----- Class: bird           AP:   0.8137 -----
----- Class: boat           AP:   0.6326 -----
----- Class: bottle         AP:   0.3017 -----
----- Class: bus            AP:   0.5275 -----
----- Class: car            AP:   0.7794 -----
----- Class: cat            AP:   0.7153 -----
----- Class: chair          AP:   0.5954 -----
----- Class: cow            AP:   0.2724 -----
----- Class: diningtable    AP:   0.4843 -----
----- Class: dog            AP:   0.6379 -----
----- Class: horse          AP:   0.6265 -----
----- Class: motorbike      AP:   0.7089 -----
----- Class: person         AP:   0.9019 -----
----- Class: pottedplant     AP:   0.3515 -----
----- Class: sheep          AP:   0.2086 -----
----- Class: sofa           AP:   0.4565 -----
----- Class: train          AP:   0.8223 -----
----- Class: tvmonitor      AP:   0.5432 -----
mAP: 0.5852
Avg loss: 0.13741200285799363
Evaluating classifier
Mean Precision Score for Testing on Epoch 1 is 0.5851549360899801
Starting epoch number 2
Loss for Training on Epoch 2 is 0.1226029172539711
Starting epoch number 3
Loss for Training on Epoch 3 is 0.10597190260887146
Starting epoch number 4
Loss for Training on Epoch 4 is 0.09369628131389618
Starting epoch number 5
Loss for Training on Epoch 5 is 0.08010298013687134
----- Class: aeroplane      AP:   0.8580 -----
----- Class: bicycle        AP:   0.7248 -----
----- Class: bird           AP:   0.8553 -----
----- Class: boat           AP:   0.7782 -----
----- Class: bottle         AP:   0.4114 -----
----- Class: bus            AP:   0.6188 -----
----- Class: car            AP:   0.8194 -----
----- Class: cat            AP:   0.7910 -----
----- Class: chair          AP:   0.6239 -----
----- Class: cow            AP:   0.4844 -----
```

-----	Class: diningtable	AP:	0.6018	-----
-----	Class: dog	AP:	0.7409	-----
-----	Class: horse	AP:	0.7250	-----
-----	Class: motorbike	AP:	0.8055	-----
-----	Class: person	AP:	0.9139	-----
-----	Class: pottedplant	AP:	0.4691	-----
-----	Class: sheep	AP:	0.5601	-----
-----	Class: sofa	AP:	0.5679	-----
-----	Class: train	AP:	0.8608	-----
-----	Class: tvmonitor	AP:	0.6367	-----

mAP: 0.6923

Avg loss: 0.11538420646798377

Evaluating classifier

Mean Precision Score for Testing on Epoch 5 is 0.6923415540683364

Starting epoch number 6

Loss for Training on Epoch 6 is 0.07785788178443909

Starting epoch number 7

Loss for Training on Epoch 7 is 0.07391707599163055

Starting epoch number 8

Loss for Training on Epoch 8 is 0.06834276020526886

Starting epoch number 9

Loss for Training on Epoch 9 is 0.05098617821931839

Starting epoch number 10

Loss for Training on Epoch 10 is 0.04607788100838661

-----	Class: aeroplane	AP:	0.8754	-----
-----	Class: bicycle	AP:	0.7110	-----
-----	Class: bird	AP:	0.8585	-----
-----	Class: boat	AP:	0.7523	-----
-----	Class: bottle	AP:	0.4112	-----
-----	Class: bus	AP:	0.6179	-----
-----	Class: car	AP:	0.8107	-----
-----	Class: cat	AP:	0.7785	-----
-----	Class: chair	AP:	0.6147	-----
-----	Class: cow	AP:	0.5179	-----
-----	Class: diningtable	AP:	0.6029	-----
-----	Class: dog	AP:	0.7296	-----
-----	Class: horse	AP:	0.7399	-----
-----	Class: motorbike	AP:	0.8026	-----
-----	Class: person	AP:	0.9033	-----
-----	Class: pottedplant	AP:	0.4436	-----
-----	Class: sheep	AP:	0.5588	-----
-----	Class: sofa	AP:	0.5481	-----
-----	Class: train	AP:	0.8862	-----
-----	Class: tvmonitor	AP:	0.6183	-----

mAP: 0.6891

Avg loss: 0.13423261776858686

Evaluating classifier

Mean Precision Score for Testing on Epoch 10 is 0.6890731384548342

Starting epoch number 11  
 Loss for Training on Epoch 11 is 0.043439846485853195  
 Starting epoch number 12  
 Loss for Training on Epoch 12 is 0.03338000923395157  
 Starting epoch number 13  
 Loss for Training on Epoch 13 is 0.02982497774064541  
 Starting epoch number 14  
 Loss for Training on Epoch 14 is 0.030931679531931877  
 Starting epoch number 15  
 Loss for Training on Epoch 15 is 0.028879933059215546  

-----	Class: aeroplane	AP:	0.8632	-----
-----	Class: bicycle	AP:	0.7091	-----
-----	Class: bird	AP:	0.8466	-----
-----	Class: boat	AP:	0.7231	-----
-----	Class: bottle	AP:	0.3739	-----
-----	Class: bus	AP:	0.6240	-----
-----	Class: car	AP:	0.8063	-----
-----	Class: cat	AP:	0.7535	-----
-----	Class: chair	AP:	0.6111	-----
-----	Class: cow	AP:	0.5133	-----
-----	Class: diningtable	AP:	0.5854	-----
-----	Class: dog	AP:	0.7100	-----
-----	Class: horse	AP:	0.7559	-----
-----	Class: motorbike	AP:	0.7934	-----
-----	Class: person	AP:	0.9076	-----
-----	Class: pottedplant	AP:	0.4340	-----
-----	Class: sheep	AP:	0.5684	-----
-----	Class: sofa	AP:	0.5459	-----
-----	Class: train	AP:	0.8865	-----
-----	Class: tvmonitor	AP:	0.6136	-----

 mAP: 0.6812  
 Avg loss: 0.15442244226441665  
 Evaluating classifier  
 Mean Precision Score for Testing on Epoch 15 is 0.6812323136585167  
 Starting epoch number 16  
 Loss for Training on Epoch 16 is 0.026101112365722656  
 Starting epoch number 17  
 Loss for Training on Epoch 17 is 0.01903168484568596  
 Starting epoch number 18  
 Loss for Training on Epoch 18 is 0.016157908365130424  
 Starting epoch number 19  
 Loss for Training on Epoch 19 is 0.029855649918317795  
 Starting epoch number 20  
 Loss for Training on Epoch 20 is 0.026157978922128677  

-----	Class: aeroplane	AP:	0.8720	-----
-----	Class: bicycle	AP:	0.6885	-----
-----	Class: bird	AP:	0.8077	-----
-----	Class: boat	AP:	0.7238	-----

-----	Class: bottle	AP: 0.3165	-----
-----	Class: bus	AP: 0.5907	-----
-----	Class: car	AP: 0.7979	-----
-----	Class: cat	AP: 0.7405	-----
-----	Class: chair	AP: 0.5722	-----
-----	Class: cow	AP: 0.5337	-----
-----	Class: diningtable	AP: 0.5388	-----
-----	Class: dog	AP: 0.6850	-----
-----	Class: horse	AP: 0.7453	-----
-----	Class: motorbike	AP: 0.7601	-----
-----	Class: person	AP: 0.8922	-----
-----	Class: pottedplant	AP: 0.4017	-----
-----	Class: sheep	AP: 0.5830	-----
-----	Class: sofa	AP: 0.5245	-----
-----	Class: train	AP: 0.8641	-----
-----	Class: tvmonitor	AP: 0.6079	-----


mAP: 0.6623


Avg loss: 0.20485579748364055

Evaluating classifier

Mean Precision Score for Testing on Epoch 20 is 0.6622990875393091

```
[27]: plot_losses(train_losses, val_losses, test_frequency, num_epochs)
      plot_mAP(train_mAPs, val_mAPs, test_frequency, num_epochs)
```





```
[28]: mAP_test, test_loss, test_aps = test_classifier(test_loader, classifier,
↪ criterion)
print("Test mAP: ", mAP_test)
```

-----	Class: aeroplane	AP: 0.8483	-----
-----	Class: bicycle	AP: 0.7503	-----
-----	Class: bird	AP: 0.7964	-----
-----	Class: boat	AP: 0.7894	-----
-----	Class: bottle	AP: 0.2776	-----
-----	Class: bus	AP: 0.6580	-----
-----	Class: car	AP: 0.8213	-----
-----	Class: cat	AP: 0.7488	-----
-----	Class: chair	AP: 0.5427	-----
-----	Class: cow	AP: 0.5058	-----
-----	Class: diningtable	AP: 0.5546	-----
-----	Class: dog	AP: 0.6797	-----
-----	Class: horse	AP: 0.8363	-----
-----	Class: motorbike	AP: 0.7136	-----
-----	Class: person	AP: 0.9069	-----
-----	Class: pottedplant	AP: 0.4180	-----
-----	Class: sheep	AP: 0.6622	-----
-----	Class: sofa	AP: 0.4921	-----
-----	Class: train	AP: 0.8469	-----
-----	Class: tvmonitor	AP: 0.5784	-----



mAP: 0.6714  
Avg loss: 0.19662140294909478  
Test mAP: 0.6713662652003508

[ ]: